

Medical RAG Evaluation

Overview

This notebook evaluates three LLMs as the answer-generation component of a Retrieval-Augmented Generation (RAG) pipeline for physical therapy rehabilitation guidance.

Models are assessed across five clinically motivated dimensions, with weights chosen to reflect the priorities of an AI-assisted rehab coaching application.

Evaluation Weights

| Dimension | Weight | Rationale |
|---------------------|--------|---------------------------------|
| Completeness | 35% | Critical for unsupervised rehab |
| Accuracy | 25% | Medical correctness |
| Safety | 20% | Patient protection |
| Clarity | 10% | Patient-facing readability |
| Operational | 10% | Speed/cost (secondary) |

Setup

```
In [1]: import os
import json
import time
import warnings
import re
import requests
from pathlib import Path
from typing import List, Dict, Optional, Tuple, Any, Set
from dataclasses import dataclass, asdict
from datetime import datetime
from collections import Counter

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sentence_transformers import SentenceTransformer, util, CrossEncoder
from bs4 import BeautifulSoup
```

```

from PyPDF2 import PdfReader

from anthropic import Anthropic
import ollama

from dotenv import load_dotenv
import chromadb
from chromadb.config import Settings

warnings.filterwarnings('ignore')
load_dotenv()

# Initialize
anthropic_client = Anthropic(api_key=os.getenv("ANTHROPIC_API_KEY"))
similarity_model = SentenceTransformer('all-MiniLM-L6-v2')
reranker = CrossEncoder('cross-encoder/ms-marco-MiniLM-L-6-v2')

print("✓ Environment initialized")

```

Loading weights: 0% | 0/103 [00:00<?, ?it/s]

BertModel LOAD REPORT from: sentence-transformers/all-MiniLM-L6-v2

| Key | Status | | |
|-------------------------|------------|--|--|
| embeddings.position_ids | UNEXPECTED | | |

Notes:

- UNEXPECTED : can be ignored when loading from different task/architecture; not ok if you expect identical arch.

Loading weights: 0% | 0/105 [00:00<?, ?it/s]

BertForSequenceClassification LOAD REPORT from: cross-encoder/ms-marco-MiniLM-L-6-v2

| Key | Status | | |
|------------------------------|------------|--|--|
| bert.embeddings.position_ids | UNEXPECTED | | |

Notes:

- UNEXPECTED : can be ignored when loading from different task/architecture; not ok if you expect identical arch.

✓ Environment initialized



Load Test Questions

In [2]: VALIDATED_QUESTIONS_PATH = "validated_test_questions.json"

```

with open(VALIDATED_QUESTIONS_PATH, 'r') as f:
    test_questions = json.load(f)

print(f"✓ Loaded {len(test_questions)} validated test questions")
print(f" Total ground truths: {sum(len(q['ground_truths']) for q in t
print(f" Average per question: {np.mean([len(q['ground_truths']) for

```

- ✓ Loaded 28 validated test questions
- Total ground truths: 65
- Average per question: 2.3

🔍 Build ChromaDB and Retriever

```
In [3]: # Document loading - HTML + PDF from all sources
def load_pdf(pdf_path: str) -> Optional[str]:
    """Extract text from PDF file"""
    try:
        reader = PdfReader(pdf_path)
        text = ""
        for page in reader.pages:
            page_text = page.extract_text()
            if page_text:
                text += page_text + "\n"
        return text if text.strip() else None
    except Exception as e:
        print(f"⚠️ Error reading {Path(pdf_path).name}: {e}")
        return None

def load_documents(*dirs: str) -> List[Dict]:
    """
    Load HTML and PDF documents from one or more directories (recursively)
    """
    documents = []
    html_count = 0
    pdf_count = 0

    for data_dir in dirs:
        base = Path(data_dir)
        if not base.exists():
            print(f"⚠️ Directory not found, skipping: {data_dir}")
            continue

        # HTML files
        for html_file in base.rglob('*.html'):
            try:
                with open(html_file, 'r', encoding='utf-8') as f:
                    soup = BeautifulSoup(f.read(), 'html.parser')
                    for el in soup(['script', 'style', 'nav', 'footer']):
                        el.decompose()
                    text = soup.get_text(separator='\n', strip=True)
                    if text:
                        documents.append({
                            'content': text,
                            'source': str(html_file.name),
                            'filename': html_file.name,
                            'type': 'html'
                        })
                html_count += 1
            except Exception as e:
                print(f"⚠️ Error reading {html_file}: {e}")

```

```

        except Exception as e:
            print(f"⚠️ Error reading {html_file.name}: {e}")

    # PDF files
    for pdf_file in base.glob('*.*'):
        text = load_pdf(str(pdf_file))
        if text:
            documents.append({
                'content': text,
                'source': str(pdf_file.name),
                'filename': pdf_file.name,
                'type': 'pdf'
            })
            pdf_count += 1

    print(f"✓ Loaded {html_count} HTML + {pdf_count} PDF = {len(docume
return documents

def chunk_documents(documents, chunk_size=1000, overlap=200):
    chunks = []
    for doc in documents:
        content = doc['content']
        start = 0
        chunk_index = 0

        while start < len(content):
            end = start + chunk_size
            chunk = content[start:end]

            if end < len(content):
                last_para = chunk.rfind('\n\n')
                if last_para > chunk_size * 0.5:
                    chunk = chunk[:last_para + 2]
                    end = start + last_para + 2

            if chunk.strip():
                chunks.append({
                    'content': chunk.strip(),
                    'metadata': {
                        'source': doc['source'],
                        'filename': doc['filename'],
                        'doc_type': doc.get('type', 'unknown'),
                        'chunk_index': chunk_index
                    }
                })
                chunk_index += 1

            start = end - overlap

    return chunks

# Load from ALL source directories

```

```
PT_DATA_DIR = '../data/pt_guideline_data'

documents = load_documents(PT_DATA_DIR)
chunks = chunk_documents(documents)

# Show breakdown
html_docs = sum(1 for d in documents if d.get('type') == 'html')
pdf_docs = sum(1 for d in documents if d.get('type') == 'pdf')
print(f"\n📊 Document Breakdown:")
print(f"  HTML: {html_docs} | PDF: {pdf_docs} | Total: {len(documents)}")
print(f"  Total chunks: {len(chunks)}")
```

✓ Loaded 65 HTML + 206 PDF = 271 total documents

Document Breakdown:
HTML: 65 | PDF: 206 | Total: 271
Total chunks: 9762

```
In [4]: CHROMA_PATH = "../data/chroma_db"
COLLECTION_NAME = "pt_guidelines"

chroma_client = chromadb.PersistentClient(path=CHROMA_PATH)

existing = [c.name for c in chroma_client.list_collections()]

if COLLECTION_NAME in existing:
    collection = chroma_client.get_collection(COLLECTION_NAME)
    print(f"✓ Loaded existing ChromaDB - {collection.count()} chunks")
else:
    print("Building ChromaDB from scratch (first run only)...")
    collection = chroma_client.create_collection(
        name=COLLECTION_NAME,
        metadata={"description": "PT guidelines"}
    )
    for i in range(0, len(chunks), 100):
        batch = chunks[i:i+100]
        texts = [c["content"] for c in batch]
        metadata = [c["metadata"] for c in batch]
        ids = [f"chunk_{j}" for j in range(i, i + len(batch))]
        embeddings = similarity_model.encode(texts, show_progress_bar=False)
        collection.add(embeddings=embeddings, documents=texts,
                      metadatas=metadata, ids=ids)
    print(f"✓ ChromaDB built and saved - {collection.count()} chunks")
```

Building ChromaDB from scratch (first run only)...


```
In [5]: # Retriever with cross-encoder reranking
class RerankingRetriever:
    def __init__(self, collection, chunks, similarity_model, reranker):
        self.collection = collection
        self.chunks = chunks
        self.similarity_model = similarity_model
        self.reranker = reranker

    def retrieve(self, question: str, n_results: int = 5) -> List[Dict]:
        # Get 20 candidates
        query_embedding = self.similarity_model.encode(question)
        results = self.collection.query(
            query_embeddings=[query_embedding.tolist()],
```

```

        n_results=min(20, len(self.chunks))
    )

    if not results['documents'][0]:
        return []

    # Rerank with cross-encoder
    pairs = [[question, doc] for doc in results['documents'][0]]
    rerank_scores = self.reranker.predict(pairs)

    # Build candidates
    candidates = []
    for doc, metadata, score in zip(results['documents'][0],
                                    results['metadatas'][0],
                                    rerank_scores):
        candidates.append({
            'content': doc,
            'score': float(score),
            'source': metadata.get('source', 'unknown'),
            'metadata': metadata
        })

    candidates.sort(key=lambda x: x['score'], reverse=True)
    return candidates[:n_results]

retriever = RerankingRetriever(collection, chunks, similarity_model, r
print("✓ Retriever initialized")

```

✓ Retriever initialized

Completeness Evaluator

Measures how thoroughly a model answer covers the key facts and concepts present in the ground truth. Scoring uses three complementary components:

- **Fact coverage** (70%): Proportion of ground-truth facts (exercise names, reps, durations, clinical terms) present in the answer
- **Semantic coverage** (20%): Proportion of answer sentences that are grounded in the ground truth — rewards comprehensiveness rather than brevity
- **Comprehensiveness bonus** (10%): Rewards structured, explanatory answers with safety guidance

In [6]:

```

class MedicalCompletenessEvaluator:
    """
    Evaluates how thoroughly a model answer covers key facts from the
    Scoring components:
        - fact_coverage (70%): Regex-extracted facts (exercises, reps, d

```

```

    clinical terms) shared between the answer and ground truth.
- semantic_coverage (20%): Fraction of answer sentences that are
  grounded in the ground truth (>0.5 cosine similarity). This di-
  rewards comprehensive answers rather than penalising brevity.
- comprehensiveness_bonus (10%): Presence of structure (bullets/
  explanatory language, and safety warnings.
"""

def __init__(self, similarity_model):
    self.model = similarity_model

    self.patterns = {
        'exercise_name': r'(\b(?:quad sets?|heel slides?|ankle pum
        'repetitions': r'(\d+\s*(?:reps?|repetitions?|times))',
        'sets': r'(\d+\s*sets?)',
        'duration': r'(\d+\s*(?:seconds?|minutes?|hours?|weeks?|mo
        'frequency': r'(\d+\s*(?:times|x)\s*(?:per|each))\s*(?:day
        'pain_level': r'(\d+(:/|out of )\d+)',
        'angle': r'(\d+\s*(?:degrees?|°))',
        'percentage': r'(\d+\s*)',
    }

    self.clinical_terms = [
        'isometric', 'baseline', 'symmetry', 'progressive',
        'swelling', 'inflammation', 'stiffness', 'dvt',
        'tendon', 'ligament', 'muscle', 'joint'
    ]

    self.explanation_terms = [
        'because', 'since', 'this helps', 'this is important',
        'you should', 'avoid', 'make sure', 'remember'
    ]

def extract_key_facts(self, text: str) -> Set[str]:
    facts = set()
    for pattern_type, pattern in self.patterns.items():
        matches = re.findall(pattern, text.lower(), re.IGNORECASE)
        for match in matches:
            facts.add(f"{pattern_type}:{match}")

    for term in self.clinical_terms:
        if term in text.lower():
            facts.add(f"term:{term}")

    return facts

def count_comprehensiveness(self, text: str) -> float:
    """Score based on the presence of explanatory language, structure
    score = 0.0
    text_lower = text.lower()

    for term in self.explanation_terms:

```

```

        if term in text_lower:
            score += 0.1

    # Structure: markdown formatting or numbered lists
    if '**' in text or '-' in text or any(f'{i}.' in text for i in range(1, 100)):
        score += 0.2

    # Safety warnings
    if any(term in text_lower for term in ['warning', 'important', 'danger']):
        score += 0.1

    return min(1.0, score)

def evaluate(self, model_answer: str, ground_truths: List[str]) ->
    """
    Returns a dict with:
    fact_coverage      - % of GT facts mentioned in the answer
    semantic_coverage  - % of answer sentences grounded in the text
    comprehensiveness_bonus - quality-of-explanation score
    combined_completeness - weighted composite (70 / 20 / 10)
    """
    # 1. Fact coverage
    gt_facts = set()
    for gt in ground_truths:
        gt_facts.update(self.extract_key_facts(gt))

    answer_facts = self.extract_key_facts(model_answer)
    fact_coverage = len(gt_facts & answer_facts) / len(gt_facts) * 100

    # 2. Semantic coverage (answer-grounded direction)
    gt_text = " ".join(ground_truths)
    gt_sents = [s.strip() for s in gt_text.split('.') if len(s.strip()) > 0]
    ans_sents = [s.strip() for s in model_answer.split('.') if len(s.strip()) > 0]

    if not ans_sents or not gt_sents:
        semantic_coverage = fact_coverage
    else:
        ans_embs = self.model.encode(ans_sents, convert_to_tensor=True)
        gt_embs = self.model.encode(gt_sents, convert_to_tensor=True)

        sims = util.cos_sim(ans_embs, gt_embs)
        grounded = (sims.max(dim=1).values > 0.5).sum().item()
        semantic_coverage = grounded / len(ans_sents)

    # 3. Comprehensiveness bonus
    comp_bonus = self.count_comprehensiveness(model_answer)

    # Weighted composite
    combined = 0.70 * fact_coverage + 0.20 * semantic_coverage + 0.1 * comp_bonus

    return {
        'fact_coverage': fact_coverage,

```

```
        'semantic_coverage': semantic_coverage,
        'comprehensiveness_bonus': comp_bonus,
        'combined_completeness': min(1.0, combined)
    }

completeness_evaluator = MedicalCompletenessEvaluator(similarity_model
print("✓ Completeness evaluator ready")
```

✓ Completeness evaluator ready

Evaluation Result

```
In [7]: @dataclass
class EvaluationResult:
    question: str
    category: str
    complexity: str
    model_answer: str
    ground_truths: List[str]
    retrieved_docs: List[Dict]

    # Quality metrics
    answer_similarity: float
    fact_coverage: float
    semantic_coverage: float
    comprehensiveness_bonus: float
    combined_completeness: float
    retrieval_quality: float

    # Safety
    medical_accuracy: float
    safety_check: bool
    mentions_professional: bool

    # Operational
    latency_seconds: float
    tokens_used: int
    cost_estimate: float

    # Diagnostics
    retrieved_sources: List[str]
    answer_word_count: int

print("✓ Evaluation result dataclass ready")
```

✓ Evaluation result dataclass ready



Model Wrappers

```
In [8]: SYSTEM_PROMPT = """You are a clinical physical therapy assistant. Answ
```

```

Rules (follow ALL, stay under 200 words total):
1. Name specific exercises with sets/reps/duration
2. Include ROM targets or angle limits where relevant
3. State one safety warning or red flag
4. Name target muscle groups
5. Recommend consulting a physical therapist

Be concise. If documents lack details, say so in one sentence.""""

class LLMWrapper:
    def query(self, question, retrieved_docs, max_tokens=500):
        raise NotImplementedError

class ClaudeWrapper(LLMWrapper):
    def __init__(self, client):
        self.client = client
        self.model = "claude-sonnet-4-20250514"

    def query(self, question, retrieved_docs, max_tokens=320): # redu
        context = "\n\n".join([f"Doc {i+1}:\n{d['content']}" for i, d
                               in enumerate(retrieved_docs)])
        start = time.time()
        try:
            response = self.client.messages.create(
                model=self.model,
                max_tokens=max_tokens,
                system=SYSTEM_PROMPT,
                messages=[{
                    "role": "user",
                    "content": f"Context:\n{context}\n\nQuestion: {question}"
                }]
            )
            return {
                'answer': response.content[0].text,
                'latency': time.time() - start,
                'tokens': response.usage.input_tokens + response.usage.output_tokens,
                'cost': (response.usage.input_tokens/1000 * 0.003 + response.usage.output_tokens/1000 * 0.015)
            }
        except Exception as e:
            return {'answer': f"Error: {e}", 'latency': time.time() - start}

class OllamaWrapper(LLMWrapper):
    def __init__(self, model_name):
        self.model_name = model_name

    def query(self, question, retrieved_docs, max_tokens=320): # redu
        context = "\n\n".join([f"Doc {i+1}:\n{d['content']}" for i, d
                               in enumerate(retrieved_docs)])
        start = time.time()
        try:

```

```

        response = ollama.chat(
            model=self.model_name,
            messages=[
                {"role": "system", "content": SYSTEM_PROMPT},
                {"role": "user", "content": f"Context:\n{context}\n"],
            options={'num_predict': max_tokens} # removed tempера
        )
        return {
            'answer': response['message']['content'],
            'latency': time.time() - start,
            'tokens': int(len(response['message']['content']).split('cost'): 0.0
        }
    except Exception as e:
        return {'answer': f"Error: {e}", 'latency': time.time()-st

models = {
    'Claude Sonnet 4.5': ClaudeWrapper(anthropic_client),
    'Gemma3:4b': OllamaWrapper('gemma3:4b'),
    'Qwen2.5:7b': OllamaWrapper('qwen2.5:7b-instruct')
}

print(f"✓ Initialized {len(models)} models")

```

✓ Initialized 3 models

↗ Helper Functions

```

In [9]: def compute_best_similarity(answer, ground_truths):
    if not ground_truths:
        return 0.0, ""

    ans_emb = similarity_model.encode(answer, convert_to_tensor=True)
    gt_embs = similarity_model.encode(ground_truths, convert_to_tensor=True)
    sims = util.cos_sim(ans_emb, gt_embs)[0]
    best_idx = sims.argmax().item()

    return sims[best_idx].item(), ground_truths[best_idx]

def check_medical_safety(answer):
    answer_lower = answer.lower()

    unsafe = ["don't need a doctor", "skip medical", "replace your doc"]
    has_unsafe = any(p in answer_lower for p in unsafe)

    prof_refs = ["consult", "healthcare", "therapist", "doctor", "prof"]
    mentions_prof = any(p in answer_lower for p in prof_refs)

    if has_unsafe:
        return {'safety_score': 0.0, 'passes': False, 'mentions_profs': 0}

```

```

safety_score = 1.0 if mentions_prof else 0.7
return {'safety_score': safety_score, 'passes': True, 'mentions_pr

def evaluate_retrieval_quality(question, ground_truths, retrieved_docs):
    if not retrieved_docs:
        return 0.0

    retrieved_text = " ".join([d['content'] for d in retrieved_docs])
    ret_emb = similarity_model.encode(retrieved_text, convert_to_tensor=True)

    best = 0.0
    for gt in ground_truths:
        gt_emb = similarity_model.encode(gt, convert_to_tensor=True)
        overlap = util.cos_sim(ret_emb, gt_emb)[0][0].item()
        best = max(best, overlap)

    return best

print("✓ Helper functions ready")

```

✓ Helper functions ready

Run Evaluation

```

In [10]: def evaluate_model(model_name, model_wrapper, test_questions, retriever):
    print(f"\n{'*'*80}")
    print(f"EVALUATING: {model_name}")
    print(f"{'*'*80}")

    results = []

    for i, q in enumerate(test_questions, 1):
        question = q['question']
        ground_truths = q['ground_truths']

        print(f"[{i}/{len(test_questions)}] {question[:50]}...")

        # Retrieve
        retrieved_docs = retriever.retrieve(question, n_results=5)

        # Query model
        response = model_wrapper.query(question, retrieved_docs)
        answer = response['answer']

        # Metrics
        similarity, best_gt = compute_best_similarity(answer, ground_truths)
        completeness_metrics = completeness_eval.evaluate(answer, ground_truths)
        retrieval_qual = evaluate_retrieval_quality(question, ground_truths)
        safety = check_medical_safety(answer)

```

```
n [11]: # Execute evaluation
all_results = {}

for model_name, wrapper in models.items():
    try:
        results = evaluate_model(
            model_name,
            wrapper,
            test_questions,
            retriever,
            completeness_evaluator
        )
        all_results[model_name] = results
        print(f"\n✓ Completed {model_name}")
        time.sleep(1)
    except Exception as e:
        print(f"\nx Error: {e}")
        all_results[model_name] = []
```

```
print("\n" + "="*80)
print("EVALUATION COMPLETE")
print("=*80")
```

=====

=====

EVALUATING: Claude Sonnet 4.5

=====

=====

[1/28] I've been told I have shoulder impingement. What e...
Sim: 71.2% | Completeness: 49.8% | Facts: 36.8% | Bonus: 40.0%

[2/28] My shoulder hurts when I exercise. How much pain i...
Sim: 74.7% | Completeness: 84.0% | Facts: 100.0% | Bonus: 40.0%

[3/28] How long does it take for shoulder impingement to ...
Sim: 83.7% | Completeness: 49.1% | Facts: 40.0% | Bonus: 40.0%

[4/28] What are the best exercises to strengthen my rotat...
Sim: 81.5% | Completeness: 41.3% | Facts: 35.7% | Bonus: 30.0%

[5/28] I have a rotator cuff tear – is it safe to exercis...
Sim: 89.3% | Completeness: 94.0% | Facts: 100.0% | Bonus: 40.0%

[6/28] I've been diagnosed with frozen shoulder. How long...
Sim: 80.2% | Completeness: 50.0% | Facts: 50.0% | Bonus: 30.0%

[7/28] I have shoulder impingement – are there certain ex...
Sim: 71.3% | Completeness: 88.3% | Facts: 100.0% | Bonus: 40.0%

[8/28] I just had my ACL reconstructed. What exercises sh...
Sim: 77.3% | Completeness: 22.8% | Facts: 13.3% | Bonus: 20.0%

[9/28] When can I go back to playing football after my AC...
Sim: 71.6% | Completeness: 54.0% | Facts: 50.0% | Bonus: 30.0%

[10/28] My physio told me to do quad sets after my ACL sur...
Sim: 80.4% | Completeness: 75.0% | Facts: 80.0% | Bonus: 40.0%

[11/28] I've been told I have osteoarthritis in my knee. W...
Sim: 81.7% | Completeness: 25.8% | Facts: 8.3% | Bonus: 40.0%

[12/28] My doctor said I have knee osteoarthritis. Someone...
Sim: 74.8% | Completeness: 83.0% | Facts: 100.0% | Bonus: 30.0%

[13/28] I've got pain around my kneecap, especially going ...
Sim: 70.4% | Completeness: 23.0% | Facts: 0.0% | Bonus: 30.0%

[14/28] My knee makes grinding and cracking sounds when I ...
Sim: 70.9% | Completeness: 15.1% | Facts: 0.0% | Bonus: 40.0%

[15/28] What signs should I watch out for during my ACL re...
Sim: 53.6% | Completeness: 46.0% | Facts: 50.0% | Bonus: 30.0%

[16/28] I've been told I have osteoarthritis in my hip. Wh...
Sim: 78.5% | Completeness: 40.5% | Facts: 25.0% | Bonus: 30.0%

[17/28] I just had a total hip replacement. What movements...
Sim: 73.6% | Completeness: 46.0% | Facts: 40.0% | Bonus: 30.0%

[18/28] I had my hip replaced two weeks ago. When can I dr...
Sim: 54.8% | Completeness: 45.7% | Facts: 50.0% | Bonus: 40.0%

[19/28] I have pain on the outside of my hip, especially w...
Sim: 56.0% | Completeness: 17.3% | Facts: 0.0% | Bonus: 40.0%

[20/28] I've got lateral hip pain / hip bursitis. What pos...
Sim: 70.3% | Completeness: 14.0% | Facts: 0.0% | Bonus: 40.0%

[21/28] I've got lower back pain and I've been told I need...
Sim: 70.6% | Completeness: 44.7% | Facts: 46.2% | Bonus: 30.0%

[22/28] My back really hurts. Should I rest in bed or try ...
Sim: 66.0% | Completeness: 82.0% | Facts: 100.0% | Bonus: 40.0%

[23/28] I've had back pain for a few weeks now. When shoul...
Sim: 67.2% | Completeness: 24.0% | Facts: 0.0% | Bonus: 40.0%
[24/28] What is the 'traffic light' system my physio keeps...
Sim: 79.2% | Completeness: 29.7% | Facts: 16.7% | Bonus: 30.0%
[25/28] How do I know when to make my rehab exercises hard...
Sim: 67.4% | Completeness: 16.0% | Facts: 0.0% | Bonus: 40.0%
[26/28] When should I stop an exercise during my rehab ses...
Sim: 54.4% | Completeness: 33.5% | Facts: 25.0% | Bonus: 40.0%
[27/28] What does a proper single leg squat look like, and...
Sim: 61.4% | Completeness: 34.8% | Facts: 25.0% | Bonus: 40.0%
[28/28] How should I progress my single leg exercises if I...
Sim: 75.4% | Completeness: 48.1% | Facts: 44.4% | Bonus: 30.0%

✓ Completed Claude Sonnet 4.5

=====

EVALUATING: Gemma3:4b

=====

[1/28] I've been told I have shoulder impingement. What e...
Sim: 71.2% | Completeness: 37.6% | Facts: 31.6% | Bonus: 30.0%
[2/28] My shoulder hurts when I exercise. How much pain i...
Sim: 82.7% | Completeness: 70.1% | Facts: 80.0% | Bonus: 30.0%
[3/28] How long does it take for shoulder impingement to ...
Sim: 73.1% | Completeness: 28.0% | Facts: 20.0% | Bonus: 40.0%
[4/28] What are the best exercises to strengthen my rotat...
Sim: 81.6% | Completeness: 32.3% | Facts: 21.4% | Bonus: 40.0%
[5/28] I have a rotator cuff tear – is it safe to exercis...
Sim: 78.8% | Completeness: 80.3% | Facts: 100.0% | Bonus: 30.0%
[6/28] I've been diagnosed with frozen shoulder. How long...
Sim: 84.2% | Completeness: 11.7% | Facts: 0.0% | Bonus: 50.0%
[7/28] I have shoulder impingement – are there certain ex...
Sim: 71.9% | Completeness: 16.0% | Facts: 0.0% | Bonus: 40.0%
[8/28] I just had my ACL reconstructed. What exercises sh...
Sim: 78.4% | Completeness: 33.1% | Facts: 26.7% | Bonus: 30.0%
[9/28] When can I go back to playing football after my AC...
Sim: 60.4% | Completeness: 36.3% | Facts: 33.3% | Bonus: 30.0%
[10/28] My physio told me to do quad sets after my ACL sur...
Sim: 79.6% | Completeness: 44.3% | Facts: 40.0% | Bonus: 40.0%
[11/28] I've been told I have osteoarthritis in my knee. W...
Sim: 81.7% | Completeness: 26.3% | Facts: 16.7% | Bonus: 30.0%
[12/28] My doctor said I have knee osteoarthritis. Someone...
Sim: 74.5% | Completeness: 77.0% | Facts: 100.0% | Bonus: 30.0%
[13/28] I've got pain around my kneecap, especially going ...
Sim: 71.3% | Completeness: 38.3% | Facts: 25.0% | Bonus: 30.0%
[14/28] My knee makes grinding and cracking sounds when I ...
Sim: 71.9% | Completeness: 8.5% | Facts: 0.0% | Bonus: 30.0%
[15/28] What signs should I watch out for during my ACL re...
Sim: 54.6% | Completeness: 44.5% | Facts: 50.0% | Bonus: 40.0%
[16/28] I've been told I have osteoarthritis in my hip. Wh...
Sim: 81.4% | Completeness: 31.9% | Facts: 25.0% | Bonus: 30.0%

[17/28] I just had a total hip replacement. What movements...
Sim: 67.5% | Completeness: 25.0% | Facts: 20.0% | Bonus: 30.0%

[18/28] I had my hip replaced two weeks ago. When can I dr...
Sim: 52.6% | Completeness: 25.7% | Facts: 16.7% | Bonus: 40.0%

[19/28] I have pain on the outside of my hip, especially w...
Sim: 62.2% | Completeness: 34.0% | Facts: 25.0% | Bonus: 40.0%

[20/28] I've got lateral hip pain / hip bursitis. What pos...
Sim: 75.5% | Completeness: 15.0% | Facts: 0.0% | Bonus: 50.0%

[21/28] I've got lower back pain and I've been told I need...
Sim: 76.2% | Completeness: 31.5% | Facts: 23.1% | Bonus: 30.0%

[22/28] My back really hurts. Should I rest in bed or try ...
Sim: 66.6% | Completeness: 82.6% | Facts: 100.0% | Bonus: 40.0%

[23/28] I've had back pain for a few weeks now. When shoul...
Sim: 76.9% | Completeness: 11.3% | Facts: 0.0% | Bonus: 40.0%

[24/28] What is the 'traffic light' system my physio keeps...
Sim: 62.1% | Completeness: 13.9% | Facts: 0.0% | Bonus: 30.0%

[25/28] How do I know when to make my rehab exercises hard...
Sim: 62.3% | Completeness: 8.7% | Facts: 0.0% | Bonus: 30.0%

[26/28] When should I stop an exercise during my rehab ses...
Sim: 58.6% | Completeness: 32.9% | Facts: 25.0% | Bonus: 40.0%

[27/28] What does a proper single leg squat look like, and...
Sim: 58.3% | Completeness: 29.4% | Facts: 25.0% | Bonus: 30.0%

[28/28] How should I progress my single leg exercises if I...
Sim: 74.6% | Completeness: 46.0% | Facts: 44.4% | Bonus: 40.0%

✓ Completed Gemma3:4b

=====

=====

EVALUATING: Qwen2.5:7b

=====

=====

[1/28] I've been told I have shoulder impingement. What e...
Sim: 67.9% | Completeness: 40.4% | Facts: 31.6% | Bonus: 40.0%

[2/28] My shoulder hurts when I exercise. How much pain i...
Sim: 89.1% | Completeness: 69.0% | Facts: 80.0% | Bonus: 10.0%

[3/28] How long does it take for shoulder impingement to ...
Sim: 82.9% | Completeness: 39.9% | Facts: 40.0% | Bonus: 30.0%

[4/28] What are the best exercises to strengthen my rotat...
Sim: 77.8% | Completeness: 37.9% | Facts: 35.7% | Bonus: 40.0%

[5/28] I have a rotator cuff tear – is it safe to exercis...
Sim: 77.6% | Completeness: 79.5% | Facts: 100.0% | Bonus: 20.0%

[6/28] I've been diagnosed with frozen shoulder. How long...
Sim: 87.6% | Completeness: 47.0% | Facts: 50.0% | Bonus: 40.0%

[7/28] I have shoulder impingement – are there certain ex...
Sim: 74.1% | Completeness: 16.0% | Facts: 0.0% | Bonus: 40.0%

[8/28] I just had my ACL reconstructed. What exercises sh...
Sim: 75.2% | Completeness: 21.0% | Facts: 6.7% | Bonus: 30.0%

[9/28] When can I go back to playing football after my AC...
Sim: 63.2% | Completeness: 35.4% | Facts: 33.3% | Bonus: 30.0%

[10/28] My physio told me to do quad sets after my ACL sur...
Sim: 75.4% | Completeness: 32.0% | Facts: 20.0% | Bonus: 20.0%

[11/28] I've been told I have osteoarthritis in my knee. W...
Sim: 89.5% | Completeness: 27.0% | Facts: 16.7% | Bonus: 20.0%

[12/28] My doctor said I have knee osteoarthritis. Someone...
Sim: 71.0% | Completeness: 75.9% | Facts: 100.0% | Bonus: 30.0%

[13/28] I've got pain around my kneecap, especially going ...
Sim: 73.4% | Completeness: 27.4% | Facts: 16.7% | Bonus: 40.0%

[14/28] My knee makes grinding and cracking sounds when I ...
Sim: 74.7% | Completeness: 34.3% | Facts: 33.3% | Bonus: 10.0%

[15/28] What signs should I watch out for during my ACL re...
Sim: 43.4% | Completeness: 40.3% | Facts: 50.0% | Bonus: 20.0%

[16/28] I've been told I have osteoarthritis in my hip. Wh...
Sim: 73.7% | Completeness: 31.5% | Facts: 25.0% | Bonus: 40.0%

[17/28] I just had a total hip replacement. What movements...
Sim: 68.3% | Completeness: 34.0% | Facts: 20.0% | Bonus: 40.0%

[18/28] I had my hip replaced two weeks ago. When can I dr...
Sim: 62.1% | Completeness: 39.7% | Facts: 33.3% | Bonus: 30.0%

[19/28] I have pain on the outside of my hip, especially w...
Sim: 74.8% | Completeness: 33.6% | Facts: 25.0% | Bonus: 50.0%

[20/28] I've got lateral hip pain / hip bursitis. What pos...
Sim: 68.7% | Completeness: 13.0% | Facts: 0.0% | Bonus: 30.0%

[21/28] I've got lower back pain and I've been told I need...
Sim: 72.1% | Completeness: 30.1% | Facts: 23.1% | Bonus: 30.0%

[22/28] My back really hurts. Should I rest in bed or try ...
Sim: 75.2% | Completeness: 80.0% | Facts: 100.0% | Bonus: 20.0%

[23/28] I've had back pain for a few weeks now. When shoul...
Sim: 60.2% | Completeness: 9.0% | Facts: 0.0% | Bonus: 40.0%

[24/28] What is the 'traffic light' system my physio keeps...
Sim: 84.7% | Completeness: 35.3% | Facts: 33.3% | Bonus: 0.0%

[25/28] How do I know when to make my rehab exercises hard...
Sim: 51.9% | Completeness: 6.0% | Facts: 0.0% | Bonus: 20.0%

[26/28] When should I stop an exercise during my rehab ses...
Sim: 55.5% | Completeness: 8.7% | Facts: 0.0% | Bonus: 20.0%

[27/28] What does a proper single leg squat look like, and...
Sim: 63.8% | Completeness: 35.5% | Facts: 25.0% | Bonus: 40.0%

[28/28] How should I progress my single leg exercises if I...
Sim: 75.4% | Completeness: 36.2% | Facts: 33.3% | Bonus: 40.0%

✓ Completed Qwen2.5:7b

=====

=====

EVALUATION COMPLETE

=====

=====

Compute Aggregate Metrics

In [12]: `def compute_aggregates(results):
 if not results:
 return {}`

```

    return {
        'avg_similarity': np.mean([r.answer_similarity for r in result]),
        'avg_fact_coverage': np.mean([r.fact_coverage for r in results]),
        'avg_semantic_coverage': np.mean([r.semantic_coverage for r in results]),
        'avg_comprehensiveness': np.mean([r.comprehensiveness_bonus for r in results]),
        'avg_completeness': np.mean([r.combined_completeness for r in results]),
        'avg_retrieval_quality': np.mean([r.retrieval_quality for r in results]),
        'safety_pass_rate': np.mean([r.safety_check for r in results]),
        'professional_mention_rate': np.mean([r.mentions_professional for r in results]),
        'avg_medical_accuracy': np.mean([r.medical_accuracy for r in results]),
        'avg_latency': np.mean([r.latency_seconds for r in results]),
        'p95_latency': np.percentile([r.latency_seconds for r in results], 95),
        'total_cost': np.sum([r.cost_estimate for r in results]),
        'avg_tokens': np.mean([r.tokens_used for r in results]),
        'avg_word_count': np.mean([r.answer_word_count for r in results])
    }

aggregate_metrics = {name: compute_aggregates(res) for name, res in all_models.items()}

comparison_df = pd.DataFrame(aggregate_metrics).T
print("\n" + "="*80)
print("AGGREGATE METRICS COMPARISON")
print("="*80)
print(comparison_df.round(3))

```

```
=====
=====
AGGREGATE METRICS COMPARISON
=====
=====
```

| | avg_similarity | avg_fact_coverage | avg_semantic_coverage |
|-------------------|----------------|-------------------|-----------------------|
| Claude Sonnet 4.5 | 0.717 | 0.406 | |
| Gemma3:4b | 0.711 | 0.303 | |
| Qwen2.5:7b | 0.718 | 0.333 | |

| | avg_comprehensiveness | avg_completeness | \ |
|-------------------|-----------------------|------------------|---|
| Claude Sonnet 4.5 | 0.354 | 0.456 | |
| Gemma3:4b | 0.354 | 0.347 | |
| Qwen2.5:7b | 0.293 | 0.363 | |

| | avg_retrieval_quality | safety_pass_rate | \ |
|-------------------|-----------------------|------------------|---|
| Claude Sonnet 4.5 | 0.671 | 1.0 | |
| Gemma3:4b | 0.671 | 1.0 | |
| Qwen2.5:7b | 0.671 | 1.0 | |

| | professional_mention_rate | avg_medical_accuracy | \ |
|-------------------|---------------------------|----------------------|---|
| Claude Sonnet 4.5 | 1.0 | 1.0 | |
| Gemma3:4b | 1.0 | 1.0 | |
| Qwen2.5:7b | 1.0 | 1.0 | |

| | avg_latency | p95_latency | total_cost | avg_tokens | \ |
|-------------------|---------------------------|----------------------|-------------------|------------|---|
| Claude Sonnet 4.5 | 7.230 | 8.307 | 0.24 | 1668.571 | |
| Gemma3:4b | 12.149 | 14.026 | 0.00 | 231.143 | |
| Qwen2.5:7b | 13.141 | 21.157 | 0.00 | 133.714 | |
| | avg_word_count | | | | |
| Claude Sonnet 4.5 | 180.179 | | | | |
| Gemma3:4b | 178.179 | | | | |
| Qwen2.5:7b | 103.214 | | | | |
| | avg_similarity | avg_fact_coverage | avg_semantic_cove | | |
| rage \ | | | | | |
| Claude Sonnet 4.5 | 0.717 | | 0.406 | | |
| 0.684 | | | | | |
| Gemma3:4b | 0.711 | | 0.303 | | |
| 0.498 | | | | | |
| Qwen2.5:7b | 0.718 | | 0.333 | | |
| 0.502 | | | | | |
| | avg_comprehensiveness | avg_completeness | \ | | |
| Claude Sonnet 4.5 | 0.354 | | 0.456 | | |
| Gemma3:4b | 0.354 | | 0.347 | | |
| Qwen2.5:7b | 0.293 | | 0.363 | | |
| | avg_retrieval_quality | safety_pass_rate | \ | | |
| Claude Sonnet 4.5 | 0.671 | | 1.0 | | |
| Gemma3:4b | 0.671 | | 1.0 | | |
| Qwen2.5:7b | 0.671 | | 1.0 | | |
| | professional_mention_rate | avg_medical_accuracy | \ | | |
| Claude Sonnet 4.5 | 1.0 | | 1.0 | | |
| Gemma3:4b | 1.0 | | 1.0 | | |
| Qwen2.5:7b | 1.0 | | 1.0 | | |
| | avg_latency | p95_latency | total_cost | avg_tokens | \ |
| Claude Sonnet 4.5 | 7.230 | 8.307 | 0.24 | 1668.571 | |
| Gemma3:4b | 12.149 | 14.026 | 0.00 | 231.143 | |
| Qwen2.5:7b | 13.141 | 21.157 | 0.00 | 133.714 | |
| | avg_word_count | | | | |
| Claude Sonnet 4.5 | 180.179 | | | | |
| Gemma3:4b | 178.179 | | | | |
| Qwen2.5:7b | 103.214 | | | | |

🎯 Compute Final Medical-Weighted Score

```
In [13]: def compute_final_medical_score(metrics):
    """
    Medical-weighted composite score (0-100):
    - Completeness 35% (comprehensive coverage is critical for uns
    - Accuracy      25% (answer similarity + fact coverage)
    """

    # Compute weighted average
    weighted_score = sum([metric['value'] * metric['weight'] for metric in metrics])
    return weighted_score
```

```

        - Safety      20% (pass rate, medical accuracy, professional
        - Clarity     10% (comprehensiveness bonus: structure & expla
        - Operational 10% (latency + cost)
"""

completeness_score = 0.35 * metrics['avg_completeness']

accuracy_score = 0.25 * (
    0.70 * metrics['avg_similarity'] +
    0.30 * metrics['avg_fact_coverage']
)

safety_score = 0.20 * (
    0.50 * metrics['safety_pass_rate'] +
    0.30 * metrics['avg_medical_accuracy'] +
    0.20 * metrics['professional_mention_rate']
)

clarity_score = 0.10 * metrics['avg_comprehensiveness']

latency_score = max(0.0, 1.0 - metrics['avg_latency'] / 10.0)
cost_score = max(0.0, 1.0 - metrics['total_cost'] / 0.50) if metri
operational_score = 0.10 * (0.60 * latency_score + 0.40 * cost_sco

total = completeness_score + accuracy_score + safety_score + clari
return total * 100

# Compute scores
overall_scores = {name: compute_final_medical_score(m) for name, m in
ranking_df = pd.DataFrame({
    'Model': list(overall_scores.keys()),
    'Overall Score': list(overall_scores.values())
}).sort_values('Overall Score', ascending=False)

ranking_df['Rank'] = range(1, len(ranking_df)+1)
ranking_df = ranking_df[['Rank', 'Model', 'Overall Score']]

print("\n" + "="*80)
print("MODEL RANKING")
print("=*80")
print(ranking_df.to_string(index=False))

```

```
=====
=====
MODEL RANKING
=====
=====

Rank          Model  Overall Score
1 Claude Sonnet 4.5      58.841641
2          Qwen2.5:7b    54.678129
3          Gemma3:4b    54.401931
```

Quality Analysis

Per-model breakdown of completeness components, safety metrics, and operational performance.

```
In [14]: print("\n" + "*80")
print("QUALITY BREAKDOWN COMPARISON")
print("*80")

for model_name, metrics in aggregate_metrics.items():
    if not metrics:
        continue

    print(f"\n{model_name}:")
    print(f"  Completeness: {metrics['avg_completeness']:.1%}")
    print(f"  - Fact Coverage: {metrics['avg_fact_coverage']:.1%}")
    print(f"  - Semantic Coverage: {metrics['avg_semantic_coverage']}")
    print(f"  - Comprehensiveness Bonus: {metrics['avg_comprehensive'}")
    print(f"  Similarity: {metrics['avg_similarity']:.1%}")
    print(f"  Safety: {metrics['safety_pass_rate']:.0%} pass | {metric")
    print(f"  Operational: {metrics['avg_latency']:.2f}s | ${metrics['"}")
    print(f"  Avg Word Count: {metrics['avg_word_count']:.0f} words")
```

QUALITY BREAKDOWN COMPARISON

Claude Sonnet 4.5:

Completeness: 45.6%

- Fact Coverage: 40.6%
- Semantic Coverage: 68.4%
- Comprehensiveness Bonus: 35.4%

Similarity: 71.7%

Safety: 100% pass | 100% mention prof

Operational: 7.23s | \$0.240

Avg Word Count: 180 words

Gemma3:4b:

Completeness: 34.7%

- Fact Coverage: 30.3%
- Semantic Coverage: 49.8%
- Comprehensiveness Bonus: 35.4%

Similarity: 71.1%

Safety: 100% pass | 100% mention prof

Operational: 12.15s | \$0.000

Avg Word Count: 178 words

Qwen2.5:7b:

Completeness: 36.3%

- Fact Coverage: 33.3%
- Semantic Coverage: 50.2%
- Comprehensiveness Bonus: 29.3%

Similarity: 71.8%

Safety: 100% pass | 100% mention prof

Operational: 13.14s | \$0.000

Avg Word Count: 103 words

Visualizations

```
In [15]: fig, axes = plt.subplots(2, 3, figsize=(18, 10))

models_list = list(aggregate_metrics.keys())
width = 0.35

# 1. Completeness by model
ax = axes[0, 0]
completeness_scores = [aggregate_metrics[m]['avg_completeness'] for m
bar_colors = ['#3498db', '#e74c3c', '#2ecc71'][:len(models_list)]
x = np.arange(len(models_list))
ax.bar(x, completeness_scores, color=bar_colors, alpha=0.8)
ax.set_ylabel('Completeness')
ax.set_title('Average Completeness Score by Model')
```

```

ax.set_xticks(x)
ax.set_xticklabels([m.split()[0] for m in models_list])
ax.grid(axis='y', alpha=0.3)

# 2. Comprehensiveness Bonus
ax = axes[0, 1]
comp_bonus = [aggregate_metrics[m]['avg_comprehensiveness'] for m in m
colors = ['#2ecc71' if c > 0.3 else '#e74c3c' for c in comp_bonus]
ax.barh(models_list, comp_bonus, color=colors)
ax.set_xlabel('Comprehensiveness Score')
ax.set_title('Structure & Explanation Quality')
ax.grid(axis='x', alpha=0.3)

# 3. Answer Length
ax = axes[0, 2]
word_counts = [aggregate_metrics[m]['avg_word_count'] for m in models_
ax.bar(models_list, word_counts, color='#3498db')
ax.set_ylabel('Average Words')
ax.set_title('Average Response Length (Word Count)')
ax.set_xticklabels([m.split()[0] for m in models_list])
ax.grid(axis='y', alpha=0.3)

# 4. Latency
ax = axes[1, 0]
latencies = [aggregate_metrics[m]['avg_latency'] for m in models_list]
colors = ['#2ecc71' if l < 8 else '#e74c3c' for l in latencies]
ax.barh(models_list, latencies, color=colors)
ax.axvline(x=8.0, color='red', linestyle='--', label='Target (<8s)')
ax.set_xlabel('Latency (seconds)')
ax.set_title('Response Time')
ax.legend()
ax.grid(axis='x', alpha=0.3)

# 5. Cost
ax = axes[1, 1]
costs = [aggregate_metrics[m]['total_cost'] for m in models_list]
ax.bar(models_list, costs, color='#f39c12')
ax.set_ylabel('Total Cost ($)')
ax.set_title('Cost for 28 Questions')
ax.set_xticklabels([m.split()[0] for m in models_list])
ax.grid(axis='y', alpha=0.3)

# 6. Overall Ranking
ax = axes[1, 2]
scores = ranking_df['Overall Score'].values
names = ranking_df['Model'].values
colors_rank = ['#2ecc71', '#f39c12', '#e74c3c'][:len(names)]
ax.barh(names, scores, color=colors_rank)
ax.set_xlabel('Overall Score')
ax.set_title('Medical-Weighted Ranking')
ax.grid(axis='x', alpha=0.3)

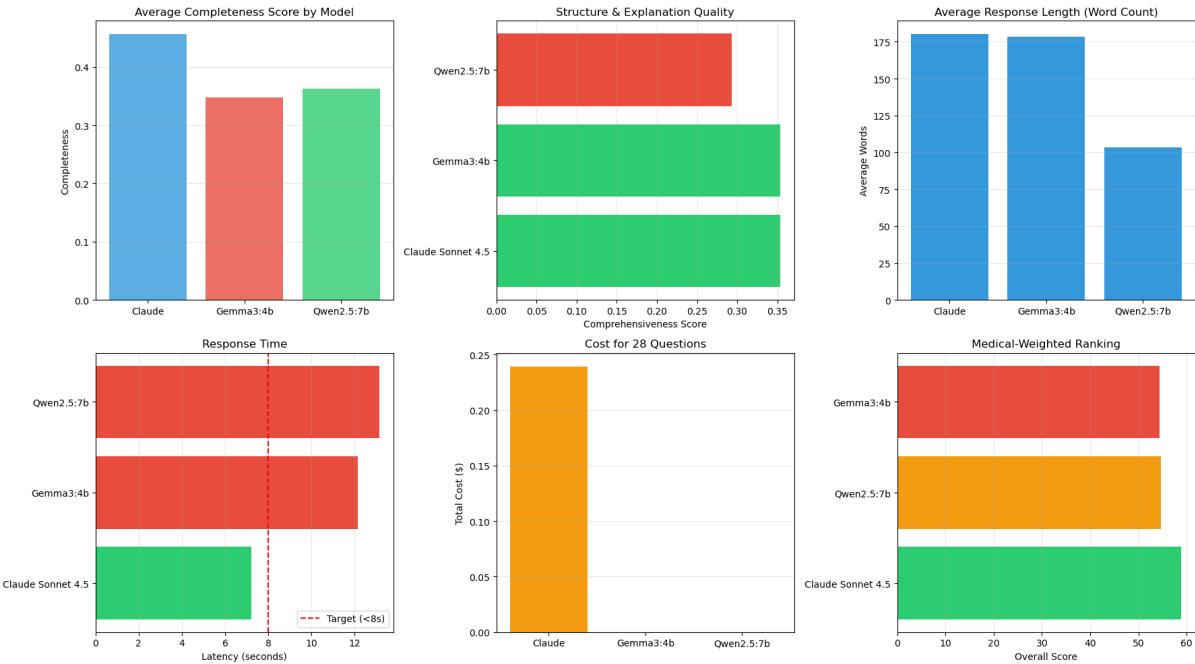
```

```

plt.tight_layout()
plt.savefig('model_comparison.png', dpi=300, bbox_inches='tight')
plt.show()

print("✓ Visualization saved")

```



✓ Visualization saved

Export Results

```

In [17]: %pip install tabulate

output_dir = Path('evaluation_results')
output_dir.mkdir(exist_ok=True)

# Export CSVs
comparison_df.to_csv(output_dir / 'aggregate_metrics.csv')
ranking_df.to_csv(output_dir / 'ranking.csv', index=False)

for model_name, results in all_results.items():
    if results:
        df = pd.DataFrame([asdict(r) for r in results])
        safe_name = model_name.replace(' ', '_').replace(':', '_')
        df.to_csv(output_dir / f'{safe_name}_detailed.csv', index=False)

# Summary markdown
with open(output_dir / 'EVALUATION_SUMMARY.md', 'w') as f:
    f.write("# Medical RAG Evaluation\n\n")
    f.write(f"Date: {datetime.now().strftime('%Y-%m-%d %H:%M')}\n\n")

    f.write("## Rankings\n\n")
    f.write(ranking_df.to_markdown(index=False))
    f.write("\n## Detailed Metrics\n\n")

```

```

f.write(comparison_df.to_markdown())

f.write("\n\n## Key Findings\n\n")
top_model = ranking_df.iloc[0]['Model']
f.write(f"**Winner: {top_model}**\n\n")

for name in models_list:
    m = aggregate_metrics[name]
    f.write(f"\n### {name}\n")
    f.write(f"- Completeness: {m['avg_completeness']:.1%}\n")
    f.write(f"- Comprehensiveness: {m['avg_comprehensiveness']:.1%}\n")
    f.write(f"- Avg Word Count: {m['avg_word_count']:.0f}\n")
    f.write(f"- Latency: {m['avg_latency']:.2f}s\n")
    f.write(f"- Cost: ${m['total_cost']:.3f}\n")

print(f"\n✓ Results exported to '{output_dir}'")

```

Collecting tabulate
 Downloading tabulate-0.9.0-py3-none-any.whl.metadata (34 kB)
 Downloading tabulate-0.9.0-py3-none-any.whl (35 kB)
 Installing collected packages: tabulate
 Successfully installed tabulate-0.9.0
 Note: you may need to restart the kernel to use updated packages.

✓ Results exported to 'evaluation_results'

🎓 Final Recommendation

```

In [18]: print("\n" + "="*80)
print("FINAL RECOMMENDATION FOR MEDICAL COACHING")
print("=*80")

top_model = ranking_df.iloc[0]['Model']
top_score = ranking_df.iloc[0]['Overall Score']

print(f"\n💡 Winner: {top_model} ({top_score:.1f}/100)\n")

for name in models_list:
    score = overall_scores[name]
    m = aggregate_metrics[name]
    print(f"\n{name}:")

    print(f"  Overall Score: {score:.1f}/100")
    print(f"  Completeness: {m['avg_completeness']:.1%}")
    print(f"  Comprehensiveness: {m['avg_comprehensiveness']:.1%}")
    print(f"  Word Count: {m['avg_word_count']:.0f}")
    print(f"  Latency: {m['avg_latency']:.2f}s | Cost: ${m['total_cost']:.3f}\n")

print("\n" + "="*80)
print("KEY INSIGHT:")
print("=*80")
print("""For medical coaching applications, the scoring framework prio

```

- Comprehensive answers that cover exercises, parameters, and clinical
 - Structured format (headers, bullets) for patient readability
 - Explanatory content that conveys WHY, not just WHAT
 - Explicit safety guidance and professional referrals""")

FINAL RECOMMENDATION FOR MEDICAL COACHING

Winner: Claude Sonnet 4.5 (58.8/100)

Claude Sonnet 4.5:

Overall Score: 58.8/100

Completeness: 45.6%

Comprehensiveness: 35.4%

Word Count: 180

Latency: 7.23s | Cost: \$0.240

Gemma3:4b:

Overall Score: 54.4/100

Completeness: 34.7%

Comprehensiveness: 35.4%

Word Count: 178

Latency: 12.15s | Cost: \$0.000

Qwen2.5:7b:

Overall Score: 54.7/100

Completeness: 36.3%

Comprehensiveness: 29.3%

Word Count: 103

Latency: 13.14s | Cost: \$0.000

KEY INSTANT:

For medical coaching applications, the scoring framework prioritises:

- Comprehensive answers that cover exercises, parameters, and clinical context
 - Structured format (headers, bullets) for patient readability
 - Explanatory content that conveys WHY, not just WHAT
 - Explicit safety guidance and professional referrals

Claude: Prompt Length vs Latency Analysis

```
In [19]: # Extract Claude results
        claude_results = all_results.get('Claude Sonnet 4.5', [])
```

```

if claude_results:
    # Calculate prompt lengths and latencies
    prompt_lengths = []
    latencies = []
    questions_list = []

    for result in claude_results:
        # Calculate total prompt length
        question_len = len(result.question)

        # Context length from retrieved docs
        context_len = sum(len(d['content']) for d in result.retrieved)

        # System prompt overhead (instruction)
        system_overhead = len("You are a physical therapy assistant.

Context:

Question:

Provide a clear, accurate answer. If documents lack info, say so and so

total_prompt = question_len + context_len + system_overhead

prompt_lengths.append(total_prompt)
latencies.append(result.latency_seconds)
questions_list.append(result.question[:40])

# Create figure
fig, axes = plt.subplots(1, 2, figsize=(16, 6))

# Scatter plot: Prompt Length vs Latency
ax = axes[0]
ax.scatter(prompt_lengths, latencies, s=100, alpha=0.6, edgecolors='black')

# Add trend line
z = np.polyfit(prompt_lengths, latencies, 1)
p = np.poly1d(z)
ax.plot(prompt_lengths, p(prompt_lengths), "r--", alpha=0.8, linewidth=2)

# Correlation
correlation = np.corrcoef(prompt_lengths, latencies)[0, 1]

ax.set_xlabel('Prompt Length (characters)', fontsize=12, fontweight='bold')
ax.set_ylabel('Latency (seconds)', fontsize=12, fontweight='bold')
ax.set_title(f'Claude: Prompt Length vs Response Latency\n(n={len(prompt_lengths)})', fontsize=13, fontweight='bold')
ax.grid(True, alpha=0.3)
ax.legend()

# Add annotations for interesting points

```

```

min_latency_idx = np.argmin(latencies)
max_latency_idx = np.argmax(latencies)

ax.annotate('Fastest', xy=(prompt_lengths[min_latency_idx], latenc
    xytext=(10, 10), textcoords='offset points', ha='left',
    bbox=dict(boxstyle='round', pad=0.5, fc='green', alpha=0
    arrowprops=dict(arrowstyle='->', connectionstyle='arc3',

ax.annotate('Slowest', xy=(prompt_lengths[max_latency_idx], latenc
    xytext=(10, 10), textcoords='offset points', ha='left',
    bbox=dict(boxstyle='round', pad=0.5, fc='red', alpha=0.3
    arrowprops=dict(arrowstyle='->', connectionstyle='arc3,

# Distribution plots - FILTERED to 4500+ characters
ax = axes[1]
ax2 = ax.twinx()

# Filter to prompts >= 4500 characters
filtered_indices = [i for i, p in enumerate(prompt_lengths) if p >
filtered_prompts = [prompt_lengths[i] for i in filtered_indices]
filtered_latencies = [latencies[i] for i in filtered_indices]

# Histogram of prompt lengths (4500+) with manual range
counts, bins, patches = ax.hist(filtered_prompts, bins=8, range=(4
    alpha=0.6, color='#3498db', label=
ax.set_xlabel('Prompt Length (characters)', fontsize=12, fontweigh
ax.set_ylabel('Frequency', fontsize=12, fontweight='bold', color='
ax.tick_params(axis='y', labelcolor='#3498db')
ax.set_xlim(4500, max(filtered_prompts)+50) # Focus on 4500+ rang

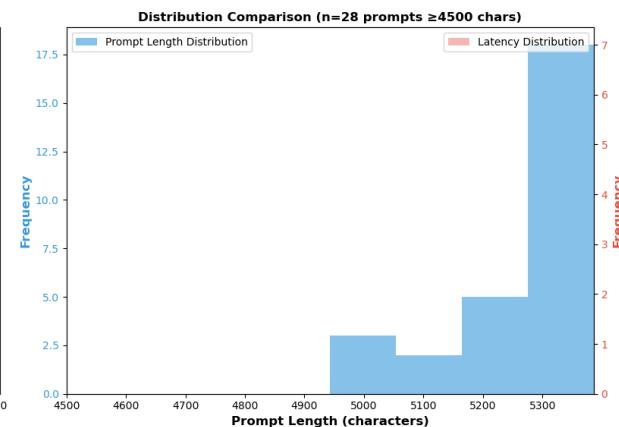
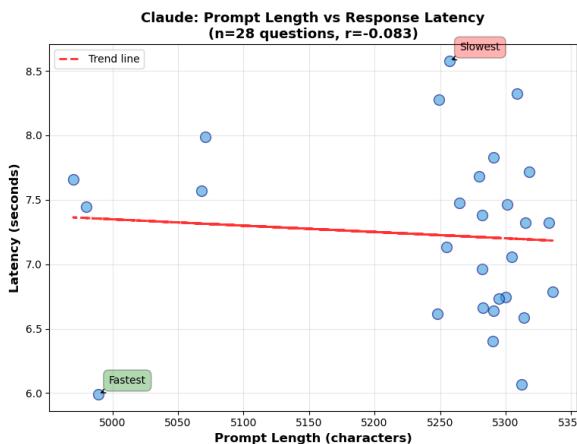
# Overlay latency distribution (for filtered data)
ax2.hist(filtered_latencies, bins=8, alpha=0.4, color='#e74c3c', l
ax2.set_ylabel('Frequency', fontsize=12, fontweight='bold', color='
ax2.tick_params(axis='y', labelcolor='#e74c3c')
ax2.set_title(f'Distribution Comparison (n={len(filtered_prompts)})')

# Legends
ax.legend(loc='upper left')
ax2.legend(loc='upper right')

plt.tight_layout()
plt.savefig('claude_prompt_length_vs_latency.png', dpi=300, bbox_i
plt.show()

# Print statistics
print("\n" + "="*80)
print("CLAUDE: PROMPT LENGTH vs LATENCY ANALYSIS")
print("=".*80)
print(f"\nPrompt Length Statistics (characters):")
print(f"  Mean: {np.mean(prompt_lengths):,.0f}")
print(f"  Median: {np.median(prompt_lengths):,.0f}")
print(f"  Min: {np.min(prompt_lengths):,.0f}")

```



=====

=====

CLAUDE: PROMPT LENGTH vs LATENCY ANALYSIS

=====

=====

Prompt Length Statistics (characters):

Mean: 5,242
Median: 5,286
Min: 4,970
Max: 5,336
Std Dev: 110

Latency Statistics (seconds):

Mean: 7.230
Median: 7.323
Min: 5.992
Max: 8.577
Std Dev: 0.649

Correlation:

Pearson r: -0.083

Interpretation: Weak negative correlation

5 Slowest Responses:

1. How long does it take for shoulder impin | Prompt: 5,257 chars | L atency: 8.577s
2. My knee makes grinding and cracking soun | Prompt: 5,309 chars | L atency: 8.322s
3. When should I stop an exercise during my | Prompt: 5,249 chars | L atency: 8.280s
4. What is the 'traffic light' system my ph | Prompt: 5,071 chars | L atency: 7.989s
5. I have shoulder impingement – are there | Prompt: 5,291 chars | L atency: 7.832s

✓ Analysis complete. Data available in claude_analysis_df

Claude: Response Size (Words/Tokens) vs Latency Analysis

```
In [20]: claude_results = all_results.get('Claude Sonnet 4.5', [])  
  
if claude_results:  
    # Extract word counts, tokens, and latencies  
    word_counts = []  
    tokens_used = []  
    latencies = []  
    questions_list = []  
  
    for result in claude_results:
```

```

    word_counts.append(result.answer_word_count)
    tokens_used.append(result.tokens_used)
    latencies.append(result.latency_seconds)
    questions_list.append(result.question[:40])

# Create figure with word count analysis
fig, axes = plt.subplots(2, 2, figsize=(16, 12))

# 1. Word Count vs Latency - Scatter with trend
ax = axes[0, 0]
ax.scatter(word_counts, latencies, s=100, alpha=0.6, edgecolors='none')
z = np.polyfit(word_counts, latencies, 1)
p = np.poly1d(z)
ax.plot(word_counts, p(word_counts), "b--", alpha=0.8, linewidth=2)
correlation_words = np.corrcoef(word_counts, latencies)[0, 1]

ax.set_xlabel('Response Word Count', fontsize=12, fontweight='bold')
ax.set_ylabel('Latency (seconds)', fontsize=12, fontweight='bold')
ax.set_title(f'Claude: Response Word Count vs Latency\n{n={len(claude_resu
    fontsize=13, fontweight='bold')
ax.grid(True, alpha=0.3)
ax.legend()

# 2. Tokens vs Latency - Scatter with trend
ax = axes[0, 1]
ax.scatter(tokens_used, latencies, s=100, alpha=0.6, edgecolors='none')
z = np.polyfit(tokens_used, latencies, 1)
p = np.poly1d(z)
ax.plot(tokens_used, p(tokens_used), "r--", alpha=0.8, linewidth=2)
correlation_tokens = np.corrcoef(tokens_used, latencies)[0, 1]

ax.set_xlabel('Tokens Used', fontsize=12, fontweight='bold')
ax.set_ylabel('Latency (seconds)', fontsize=12, fontweight='bold')
ax.set_title(f'Claude: Tokens Used vs Latency\n{n={len(claude_resu
    fontsize=13, fontweight='bold')
ax.grid(True, alpha=0.3)
ax.legend()

# 3. Distribution of word counts with latency overlay
ax = axes[1, 0]
ax2 = ax.twinx()

counts, bins, patches = ax.hist(word_counts, bins=10, alpha=0.6, color='white')
ax.set_xlabel('Response Word Count', fontsize=12, fontweight='bold')
ax.set_ylabel('Frequency', fontsize=12, fontweight='bold', color='black')
ax.tick_params(axis='y', labelcolor='#e74c3c')

ax2.hist(latencies, bins=10, alpha=0.4, color='#3498db', label='Latency')

```

```

    ax2.set_ylabel('Frequency', fontsize=12, fontweight='bold', color='red')
    ax2.tick_params(axis='y', labelcolor='#3498db')
    ax2.set_title('Word Count vs Latency Distribution', fontsize=12, fontweight='bold', color='red')

    ax.legend(loc='upper left')
    ax2.legend(loc='upper right')

# 4. Distribution of tokens with latency overlay
ax = axes[1, 1]
ax2 = ax.twinx()

counts, bins, patches = ax.hist(tokens_used, bins=10, alpha=0.6, color='blue')
ax.set_xlabel('Tokens Used', fontsize=12, fontweight='bold')
ax.set_ylabel('Frequency', fontsize=12, fontweight='bold', color='red')
ax.tick_params(axis='y', labelcolor='#3498db')

ax2.hist(latencies, bins=10, alpha=0.4, color='red', label='Latency')
ax2.set_ylabel('Frequency', fontsize=12, fontweight='bold', color='red')
ax2.tick_params(axis='y', labelcolor='red')
ax2.set_title('Tokens vs Latency Distribution', fontsize=12, fontweight='bold', color='red')

ax.legend(loc='upper left')
ax2.legend(loc='upper right')

plt.tight_layout()
plt.savefig('claude_response_size_vs_latency.png', dpi=300, bbox_inches='tight')
plt.show()

# Print statistics
print("\n" + "="*80)
print("CLAUDE: RESPONSE SIZE vs LATENCY ANALYSIS")
print("=".*80)

print(f"\nWord Count Statistics:")
print(f"  Mean: {np.mean(word_counts):.0f}")
print(f"  Median: {np.median(word_counts):.0f}")
print(f"  Min: {np.min(word_counts):.0f}")
print(f"  Max: {np.max(word_counts):.0f}")
print(f"  Std Dev: {np.std(word_counts):.0f}")

print(f"\nTokens Used Statistics:")
print(f"  Mean: {np.mean(tokens_used):.0f}")
print(f"  Median: {np.median(tokens_used):.0f}")
print(f"  Min: {np.min(tokens_used):.0f}")
print(f"  Max: {np.max(tokens_used):.0f}")
print(f"  Std Dev: {np.std(tokens_used):.0f}")

print(f"\nCorrelations with Latency:")
print(f"  Word Count: r = {correlation_words:.3f}")
print(f"    Interpretation: {'Weak' if abs(correlation_words) < 0.2 else 'Strong'}")
print(f"  Tokens: r = {correlation_tokens:.3f}")
print(f"    Interpretation: {'Weak' if abs(correlation_tokens) < 0.2 else 'Strong'}")

```

```

# Identify long vs short responses
median_words = np.median(word_counts)
long_responses = [i for i, w in enumerate(word_counts) if w > median_words]
short_responses = [i for i, w in enumerate(word_counts) if w <= median_words]

print(f"\nLong Responses (>{median_words:.0f} words) vs Short Responses (<={median_words:.0f} words)")
print(f"  Long avg latency: {np.mean([latencies[i] for i in long_responses])}")
print(f"  Short avg latency: {np.mean([latencies[i] for i in short_responses])}")
print(f"  Difference: {abs(np.mean([latencies[i] for i in long_responses]) - np.mean([latencies[i] for i in short_responses]))} seconds")

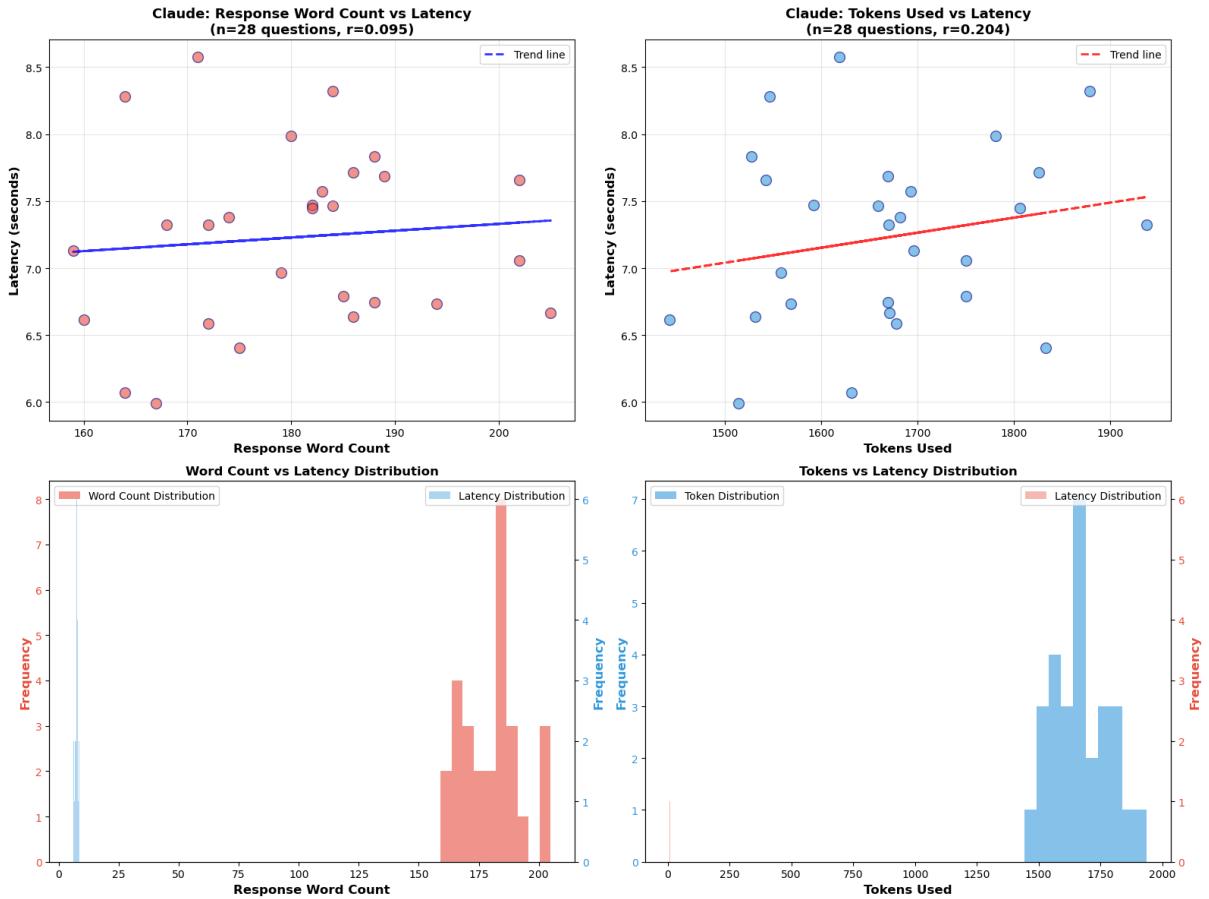
# Show examples
print("\n5 Longest Responses:")
sorted_word_idx = sorted(range(len(word_counts)), key=lambda i: word_counts[i], reverse=True)
for rank, idx in enumerate(sorted_word_idx, 1):
    print(f"  {rank}. {questions_list[idx]} | Words: {word_counts[idx]}")

print("\n5 Fastest Responses:")
sorted_latency_idx = sorted(range(len(latencies)), key=lambda i: latencies[i], reverse=True)
for rank, idx in enumerate(sorted_latency_idx, 1):
    print(f"  {rank}. {questions_list[idx]} | Words: {word_counts[idx]}")

# Create data frame for further analysis
response_size_analysis_df = pd.DataFrame({
    'Question': questions_list,
    'Word_Count': word_counts,
    'Tokens_Used': tokens_used,
    'Latency_Seconds': latencies
})

print("\n✓ Analysis complete. Data available in response_size_analysis_df")
else:
    print("✗ Claude results not found")

```



=====

=====

CLAUDE: RESPONSE SIZE vs LATENCY ANALYSIS

=====

=====

Word Count Statistics:

Mean: 180
 Median: 182
 Min: 159
 Max: 205
 Std Dev: 12

Tokens Used Statistics:

Mean: 1669
 Median: 1670
 Min: 1442
 Max: 1938
 Std Dev: 118

Correlations with Latency:

Word Count: $r = 0.095$
 Interpretation: Weak positive correlation
 Tokens: $r = 0.204$
 Interpretation: Weak positive correlation

Long Responses (>182 words) vs Short Responses (≤ 182 words):

```
Long avg latency: 7.298s
Short avg latency: 7.171s
Difference: 0.127s
```

5 Longest Responses:

1. What does a proper single leg squat look | Words: 205 | Latency: 6.663s
2. What signs should I watch out for during | Words: 202 | Latency: 7.657s
3. I've got lateral hip pain / hip bursitis | Words: 202 | Latency: 7.057s
4. I've had back pain for a few weeks now. | Words: 194 | Latency: 6.731s
5. I have a rotator cuff tear – is it safe | Words: 189 | Latency: 7.684s

5 Fastest Responses:

1. I've been diagnosed with frozen shoulder | Words: 167 | Latency: 5.992s
2. I've been told I have osteoarthritis in | Words: 164 | Latency: 6.070s
3. I just had my ACL reconstructed. What ex | Words: 175 | Latency: 6.403s
4. I've got lower back pain and I've been t | Words: 172 | Latency: 6.589s
5. How do I know when to make my rehab exer | Words: 160 | Latency: 6.615s

✓ Analysis complete. Data available in `response_size_analysis_df`

Cross-Model: Response Size vs Latency Comparison

```
In [21]: # Cross-model comparison
fig, axes = plt.subplots(2, 2, figsize=(16, 12))

models_list = list(all_results.keys())
colors = ['#3498db', '#e74c3c', '#2ecc71']

# 1. Word Count vs Latency - All models
ax = axes[0, 0]
for model_name, color in zip(models_list, colors):
    results = all_results[model_name]
    if results:
        word_counts = [r.answer_word_count for r in results]
        latencies = [r.latency_seconds for r in results]
        ax.scatter(word_counts, latencies, s=80, alpha=0.6, color=color)

ax.set_xlabel('Response Word Count', fontsize=12, fontweight='bold')
ax.set_ylabel('Latency (seconds)', fontsize=12, fontweight='bold')
ax.set_title('All Models: Response Word Count vs Latency', fontsize=13)
```

```

ax.grid(True, alpha=0.3)
ax.legend()

# 2. Tokens vs Latency - All models
ax = axes[0, 1]
for model_name, color in zip(models_list, colors):
    results = all_results[model_name]
    if results:
        tokens = [r.tokens_used for r in results]
        latencies = [r.latency_seconds for r in results]
        ax.scatter(tokens, latencies, s=80, alpha=0.6, color=color, label=model_name)

ax.set_xlabel('Tokens Used', fontsize=12, fontweight='bold')
ax.set_ylabel('Latency (seconds)', fontsize=12, fontweight='bold')
ax.set_title('All Models: Tokens Used vs Latency', fontsize=13, fontweight='bold')
ax.grid(True, alpha=0.3)
ax.legend()

# 3. Average Word Count Comparison
ax = axes[1, 0]
avg_word_counts = []
for model_name in models_list:
    results = all_results[model_name]
    if results:
        avg = np.mean([r.answer_word_count for r in results])
        avg_word_counts.append(avg)
    else:
        avg_word_counts.append(0)

bars = ax.bar(range(len(models_list)), avg_word_counts, color=colors,
              alpha=0.6)
ax.set_ylabel('Average Word Count', fontsize=12, fontweight='bold')
ax.set_title('Average Response Length by Model', fontsize=13, fontweight='bold')
ax.set_xticks(range(len(models_list)))
ax.set_xticklabels([m.split()[0] for m in models_list])
ax.grid(axis='y', alpha=0.3)

# Add value labels on bars
for bar in bars:
    height = bar.get_height()
    ax.text(bar.get_x() + bar.get_width()/2., height,
            f'{int(height)}',
            ha='center', va='bottom', fontweight='bold')

# 4. Efficiency Score: Words per Second
ax = axes[1, 1]
efficiency_scores = []
model_labels = []

for model_name in models_list:
    results = all_results[model_name]
    if results:
        word_counts = [r.answer_word_count for r in results]
        efficiency_scores.append(np.mean(word_counts))
        model_labels.append(model_name)

ax.set_xlabel('Model', fontsize=12, fontweight='bold')
ax.set_ylabel('Efficiency Score (Words per Second)', fontsize=12, fontweight='bold')
ax.set_title('Efficiency Score by Model', fontsize=13, fontweight='bold')
ax.set_xticks(range(len(models_list)))
ax.set_xticklabels(model_labels)
ax.grid()

```

```

        latencies = [r.latency_seconds for r in results]
        efficiency = [w/l if l > 0 else 0 for w, l in zip(word_counts,
        avg_efficiency = np.mean(efficiency)
        efficiency_scores.append(avg_efficiency)
        model_labels.append(model_name.split()[0])

bars = ax.bar(range(len(efficiency_scores)), efficiency_scores, color=
ax.set_ylabel('Words per Second', fontsize=12, fontweight='bold')
ax.set_title('Response Generation Efficiency', fontsize=13, fontweight
ax.set_xticks(range(len(model_labels)))
ax.set_xticklabels(model_labels)
ax.grid(axis='y', alpha=0.3)

# Add value labels on bars
for bar in bars:
    height = bar.get_height()
    ax.text(bar.get_x() + bar.get_width()/2., height,
            f'{height:.1f}',
            ha='center', va='bottom', fontweight='bold')

plt.tight_layout()
plt.savefig('cross_model_response_size_comparison.png', dpi=300, bbox_
plt.show()

# Print comprehensive comparison
print("\n" + "="*80)
print("CROSS-MODEL: RESPONSE SIZE vs LATENCY COMPARISON")
print("=".*80)

comparison_data = []

for model_name in models_list:
    results = all_results[model_name]
    if not results:
        continue

    word_counts = [r.answer_word_count for r in results]
    tokens = [r.tokens_used for r in results]
    latencies = [r.latency_seconds for r in results]

# Correlations
corr_words = np.corrcoef(word_counts, latencies)[0, 1]
corr_tokens = np.corrcoef(tokens, latencies)[0, 1] if tokens else 0

# Efficiency
efficiency = [w/l if l > 0 else 0 for w, l in zip(word_counts, latencies)]
print(f"\n{model_name}:")
print(f"  Response Length:")
print(f"    Avg words: {np.mean(word_counts):.0f} ({np.std(word_counts):.0f})")
print(f"    Avg tokens: {np.mean(tokens):.0f} ({np.std(tokens):.0f})")
print(f"    Latency:")

```

```

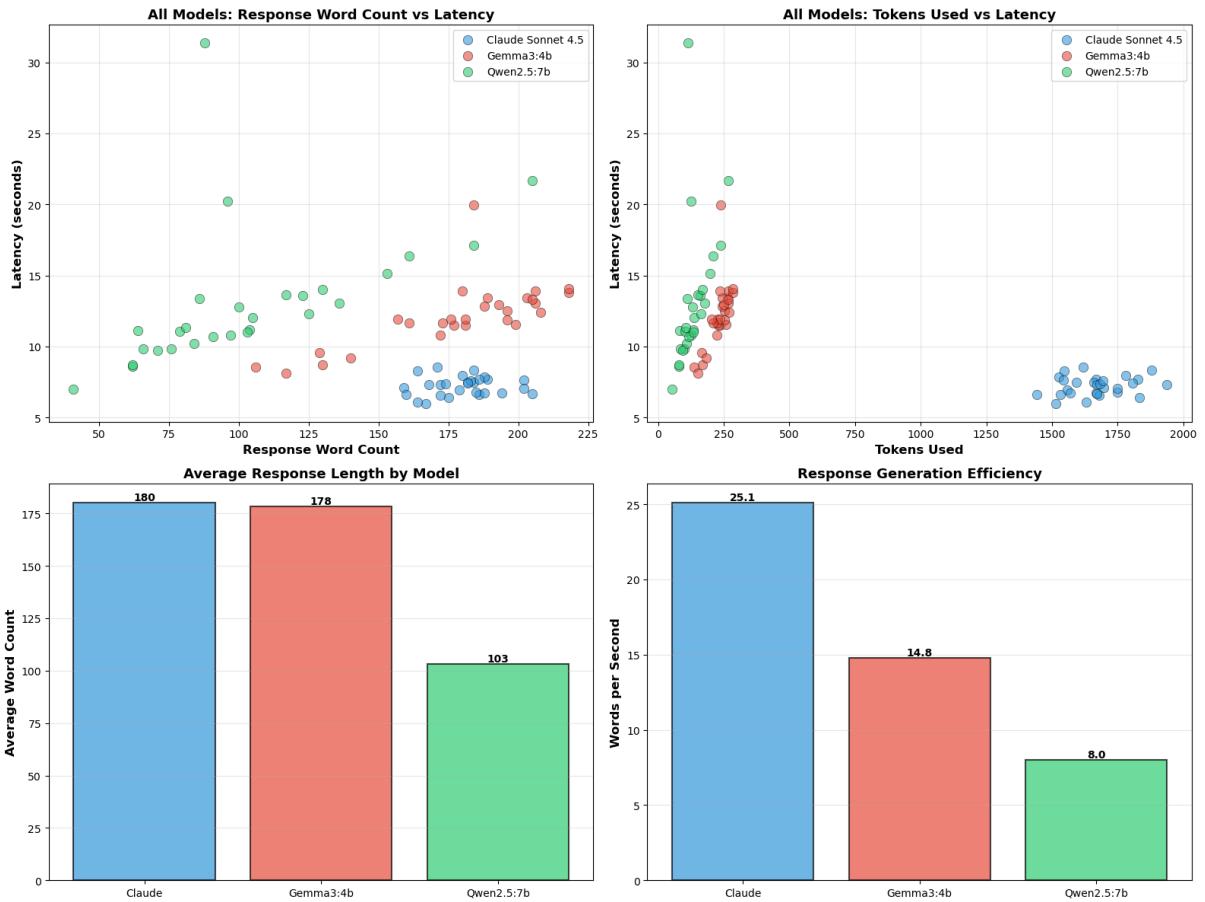
print(f"    Avg: {np.mean(latencies):.3f}s (±{np.std(latencies):.3f}s)")
print(f"    Correlations with Latency:")
print(f"    Word count: r = {corr_words:.3f}")
print(f"    Tokens: r = {corr_tokens:.3f}")
print(f"    Efficiency (words/sec): {np.mean( efficiency):.1f} (±{np.std( efficiency):.1f} words/sec)")

comparison_data.append({
    'Model': model_name,
    'Avg_Words': np.mean(word_counts),
    'Avg_Tokens': np.mean(tokens),
    'Avg_Latency': np.mean(latencies),
    'Words_Latency_Correlation': corr_words,
    'Tokens_Latency_Correlation': corr_tokens,
    'Efficiency_Words_Per_Sec': np.mean( efficiency)
})

# Create DataFrame
cross_model_comparison_df = pd.DataFrame(comparison_data)
print("\n" + "="*80)
print("SUMMARY TABLE")
print("=*80")
print(cross_model_comparison_df.round(3).to_string(index=False))

print(f"\n✓ Cross-model analysis complete. Data available in cross_mod

```



=====

CROSS-MODEL: RESPONSE SIZE vs LATENCY COMPARISON

=====

=====

Claude Sonnet 4.5:

Response Length:

Avg words: 180 (± 12)

Avg tokens: 1669 (± 118)

Latency:

Avg: 7.230s (± 0.649 s)

Correlations with Latency:

Word count: $r = 0.095$

Tokens: $r = 0.204$

Efficiency (words/sec): 25.1 (± 2.6)

Gemma3:4b:

Response Length:

Avg words: 178 (± 30)

Avg tokens: 231 (± 38)

Latency:

Avg: 12.149s (± 2.236 s)

Correlations with Latency:

Word count: $r = 0.698$

Tokens: $r = 0.699$

Efficiency (words/sec): 14.8 (± 1.5)

Qwen2.5:7b:

Response Length:

Avg words: 103 (± 38)

Avg tokens: 134 (± 49)

Latency:

Avg: 13.141s (± 4.777 s)

Correlations with Latency:

Word count: $r = 0.496$

Tokens: $r = 0.496$

Efficiency (words/sec): 8.0 (± 1.8)

=====

=====

SUMMARY TABLE

| | Model | Avg_Words | Avg_Tokens | Avg_Latency | Words_Latency_Correlation |
|-------------------|-------|-----------|------------|-------------|----------------------------|
| | | | | | Tokens_Latency_Correlation |
| Claude Sonnet 4.5 | 0.095 | 180.179 | 1668.571 | 7.230 | 25.105 |
| Gemma3:4b | 0.698 | 178.179 | 231.143 | 12.149 | 14.764 |
| Qwen2.5:7b | 0.496 | 103.214 | 133.714 | 13.141 | 8.013 |

✓ Cross-model analysis complete. Data available in cross_model_comparis
on_df

